

Assignment 1

MAD

Modelling and Analysis of Data
NDAB16012U

DEPARTMENT OF COMPUTER SCIENCE
UNIVERSITY OF COPENHAGEN

Jack Bjerregaard - 1dh587

23/11/2025

1 Partial Derivatives

a) To find the partial derivate of $f(x, y) = x^4y^3 + 7x^5 - e^{xy}$ we begin by finding the partial derivative with respect to x , where we treat y as a constant, and then do the same with respect to y :

$$f(x, y) = x^4y^3 + 7x^5 - e^{xy} \quad (1)$$

$$\frac{\partial f}{\partial x} = \frac{\partial}{\partial x} (x^4y^3 + 7x^5 - e^{xy}) \quad (2)$$

$$= 4x^3y^3 + 35x^4 - ye^{xy} \quad (3)$$

$$\frac{\partial f}{\partial y} = \frac{\partial}{\partial y} (x^4y^3 + 7x^5 - e^{xy}) \quad (4)$$

$$= 3x^4y^2 - xe^{xy} \quad (5)$$

b) Similar to how we solved the previous partial derivatives, we again attempt to find the partial derivatives of each variable, treating the other as a constant. However, to simplify it, we begin by rewriting the function with exponent $-\frac{1}{2}$ and then using the chain rule:

$$f(x, y) = \frac{1}{\sqrt{x^3 + xy + y^2}} = (x^3 + xy + y^2)^{-1/2} \quad (6)$$

$$\frac{\partial f}{\partial x} = \frac{\partial}{\partial x} (x^3 + xy + y^2)^{-1/2} \quad (7)$$

$$= -\frac{1}{2} (x^3 + xy + y^2)^{-3/2} \cdot \frac{\partial}{\partial x} (x^3 + xy + y^2) \quad (8)$$

$$= -\frac{1}{2} (x^3 + xy + y^2)^{-3/2} (3x^2 + y) \quad (9)$$

$$= -\frac{3x^2 + y}{2\sqrt{(x^3 + xy + y^2)^3}} \quad (10)$$

$$\frac{\partial f}{\partial y} = \frac{\partial}{\partial y} (x^3 + xy + y^2)^{-1/2} \quad (11)$$

$$= -\frac{1}{2} (x^3 + xy + y^2)^{-3/2} \cdot \frac{\partial}{\partial y} (x^3 + xy + y^2) \quad (12)$$

$$= -\frac{1}{2} (x^3 + xy + y^2)^{-3/2} (x + 2y) \quad (13)$$

$$= -\frac{x + 2y}{2\sqrt{(x^3 + xy + y^2)^3}} \quad (14)$$

c) To solve the last partial derivative problem, we do the same as in b), however instead of the chain rule we use the quotient rule:

$$f(x, y) = \frac{x^3 + y^2}{x + y} \quad (15)$$

$$\frac{\partial f}{\partial x} = \frac{\partial}{\partial x} \left(\frac{x^3 + y^2}{x + y} \right) \quad (16)$$

$$= \frac{(3x^2)(x + y) - (x^3 + y^2)(1)}{(x + y)^2} \quad (17)$$

$$= \frac{3x^3 + 3x^2y - x^3 - y^2}{(x + y)^2} \quad (18)$$

$$= \frac{2x^3 + 3x^2y - y^2}{(x + y)^2} \quad (19)$$

$$\frac{\partial f}{\partial y} = \frac{\partial}{\partial y} \left(\frac{x^3 + y^2}{x + y} \right) \quad (20)$$

$$= \frac{(2y)(x + y) - (x^3 + y^2)(1)}{(x + y)^2} \quad (21)$$

$$= \frac{2xy + 2y^2 - x^3 - y^2}{(x + y)^2} \quad (22)$$

$$= \frac{-x^3 + 2xy + y^2}{(x + y)^2} \quad (23)$$

2 Gradients

a) We have a function of a vector, $f(\bar{x}) = \bar{x}^T \bar{x} + c$. To compute the gradient ∇f with respect to \bar{x} we note the definition from Table 1.4 (Rogers & Girolami, 2012, p.23), where we have that the partial derivative with respect to w of $w^T w = 2w$. Thus, we have:

$$\frac{\partial}{\partial w}(w^T w) = 2w.$$

Substituting the definition for our question so $w = \bar{x}$,

$$\nabla_{\bar{x}} f = 2\bar{x}.$$

Since our constant c is independent of \bar{x} , its derivative is zero and our final gradient is:

$$\nabla f = 2\bar{x}.$$

b) We now consider the linear function of a vector

$$f(\bar{x}) = \bar{x}^T \bar{b},$$

Where we have that \bar{b} is a constant vector. To compute the gradient ∇f with respect to \bar{x} , we use the identity from table 1.4 (Rogers & Girolami, 2012, p.23), which states that

for a given vector w and a constant vector x we have:

$$\frac{\partial}{\partial w}(w^T x) = x.$$

Then we substitute for our question so $w = \bar{x}$ and $x = \bar{b}$:

$$\nabla_{\bar{x}} f = \frac{\partial}{\partial \bar{x}}(\bar{x}^T \bar{b}) = \bar{b}.$$

Thus, the gradient of f with respect to \bar{x} is

$$\nabla f = \bar{b}.$$

c) For our final function, we have a function of a quadratic term, linear term and constant:

$$f(\bar{x}) = \bar{x}^T A \bar{x} + \bar{b}^T \bar{x} + c,$$

where A is a matrix, \bar{b} is a constant vector, and c is a constant scalar. We compute the gradient ∇f with respect to \bar{x} by differentiating each term separately.

From Table 1.4 we use the identity that for a vector w and a matrix C

$$\frac{\partial}{\partial w}(w^T C w) = (C + C^T)w.$$

Taking $w = \bar{x}$ and $C = A$ gives

$$\frac{\partial}{\partial \bar{x}}(\bar{x}^T A \bar{x}) = (A + A^T)\bar{x}.$$

For the linear term, we get the same result as part (b):

$$\frac{\partial}{\partial \bar{x}}(\bar{b}^T \bar{x}) = \bar{b}.$$

The constant c does not depend on \bar{x} , so its derivative is zero:

$$\frac{\partial}{\partial \bar{x}}c = 0.$$

thus, we get the full gradient:

$$\nabla_{\bar{x}} f = (A + A^T)\bar{x} + \bar{b}.$$

So the gradient of f with respect to \bar{x} is

$$\nabla f = (A + A^T)\bar{x} + \bar{b}.$$

3 Estimating House Prices I

a) To implement the naive approach, we resorted to using the average price for all the houses in the training set. The mean was found to be $\approx \$22.02k$.

b) After implementing and using the `rmse(t, tp)` function, the computed RMSE was found to be $\approx \$9.67k$

C) Using the values computed, and `matplotlib` a 2D scatter plot was created to visualize this naive approach. This can be found in the following figure:

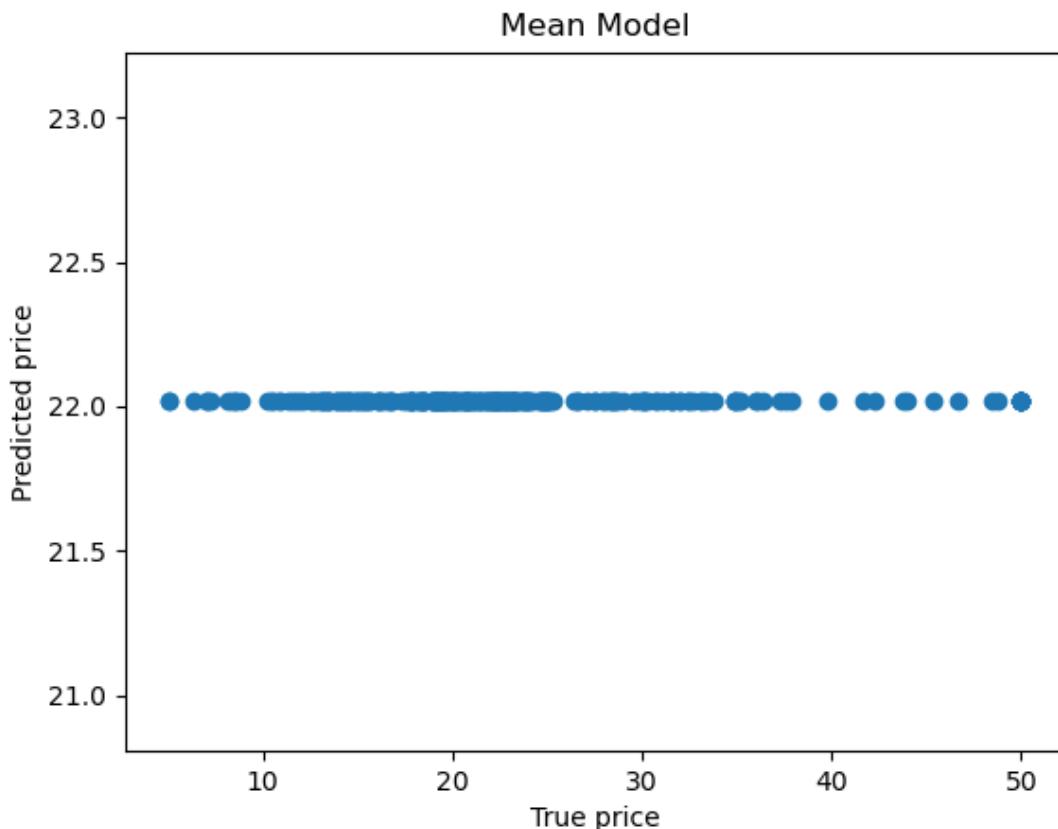


Figure 1: Estimating House Prices using Naive approach

4 Estimating House prices II

a) Our solution to a) can be found in the respective `src` files (`housing_2.py` and `linreg.py`). Our `fit` function is as follows:

```
1 def fit(self, X, t):
2     """
3         Fits the linear regression model.
4
5     Parameters
6     -----
7     X : Array of shape [n_samples, n_features]
```

```
8     t : Array of shape [n_samples, 1]
9     """
10    if X.ndim == 1:
11        X = X.reshape(-1, 1)  # converts 1d vector to matrix
12
13    ones = numpy.ones((X.shape[0], 1))
14    new_X = numpy.hstack([ones, X])  # make new matrix with 1s
15
16    # building the normal equation - Aw = b, solving for w
17    A = new_X.T @ new_X
18    b = new_X.T @ t
19    self.w = numpy.linalg.solve(A, b)
20
```

b) Using the code from a) a linear regression model was fitted on the training set only containing the first feature. From this we found the 'intercept' to be $\approx \$23.6k$ and the 'CRIM' to be ≈ -0.433 . The intercept can be interpreted as being a representation of the predicted house prices (per \$1000) when the crime rate is zero, while the coefficient (CRIM) indicates that for each unit increase in crime rate, the house price decreases by approximately \$433.

c) The computed weights are in the following table:

Index	Weight	Feature
w_0	31.39	Intercept
w_1	-0.0596	CRIM
w_2	0.0294	ZN
w_3	-0.0291	INDUS
w_4	2.293	CHAS
w_5	-17.33	NOX
w_6	3.994	RM
w_7	0.00323	AGE
w_8	-1.287	DIS
w_9	0.3548	RAD
w_{10}	-0.0156	PTRATIO
w_{11}	-0.8146	TAX
w_{12}	0.0118	B
w_{13}	-0.4649	LSTAT

Table 1: Linear regression weights for the full feature model.

d) After implementing the `predict` function, we used it to compute predictions for both models. From this we found that the RMSE for single was $\approx \$8.96k$ and the RMSE for all features was $\approx \$4.69k$. The scatter plots for both features are in the following plots:

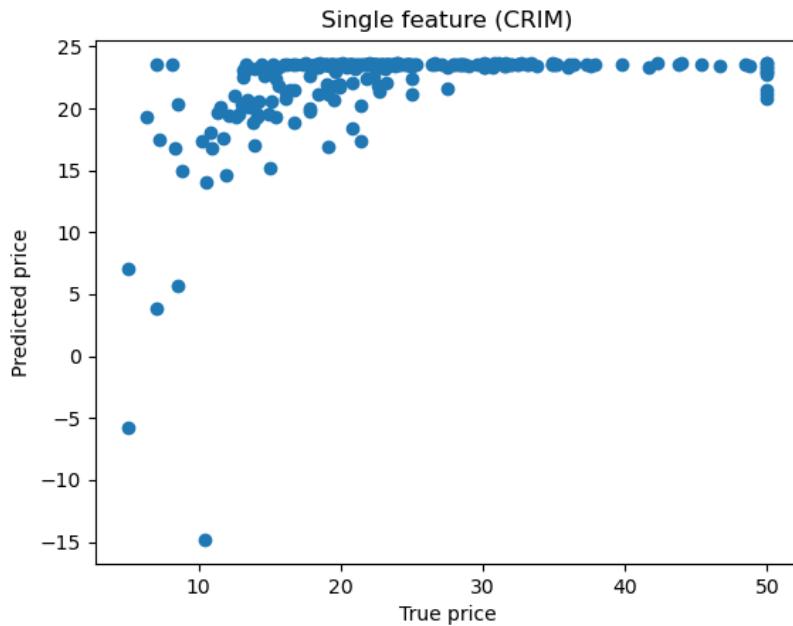


Figure 2: Scatter single

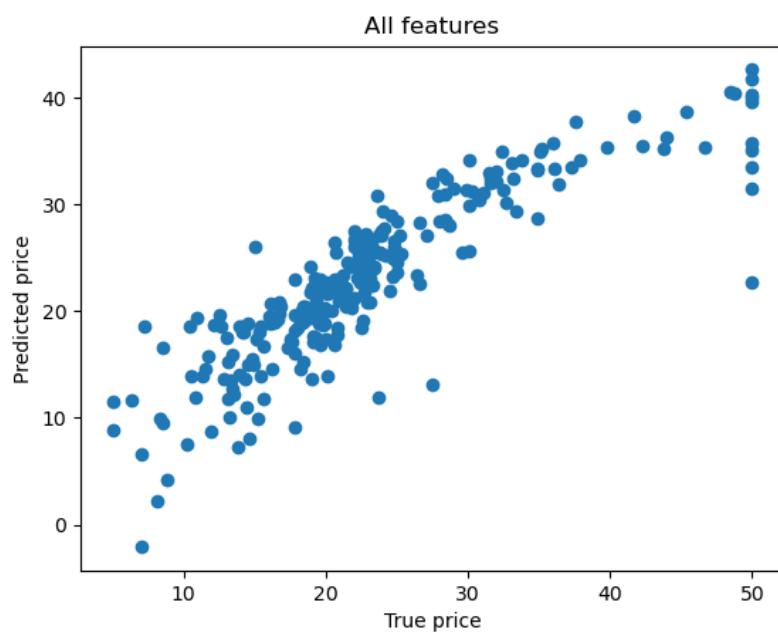


Figure 3: Scatter all

5 Total Training Loss

a) Given the total loss function:

$$\mathcal{L}(\mathbf{w}) = \sum_{n=1}^N (\mathbf{w}^T \mathbf{x}_n - t_n)^2 \quad (24)$$

we begin by rewriting it in matrix form to simplify differentiation and computation. Following section 1.3, (Rogers & Girolami), we define the data matrix X and the target vector t : We begin by rewriting it in matrix form to simplify differentiation and computation.

$$X = \begin{bmatrix} \mathbf{x}_1^T \\ \mathbf{x}_2^T \\ \vdots \\ \mathbf{x}_N^T \end{bmatrix}, \quad \mathbf{t} = \begin{bmatrix} t_1 \\ t_2 \\ \vdots \\ t_N \end{bmatrix} \quad (25)$$

Using this notation, the loss can be expressed compactly as (see Equation 1.13)(Rogers & Girolami, p.21):

$$\mathcal{L}(\mathbf{w}) = (X\mathbf{w} - \mathbf{t})^T (X\mathbf{w} - \mathbf{t}) \quad (26)$$

This expression can be expanded using the transpose of a product (Rogers & Girolami, comment 1.8, p.21):

$$\mathcal{L}(\mathbf{w}) = ((X\mathbf{w})^T - \mathbf{t}^T)(X\mathbf{w} - \mathbf{t}) \quad (27)$$

$$= (X\mathbf{w})^T X\mathbf{w} - \mathbf{t}^T X\mathbf{w} - (X\mathbf{w})^T \mathbf{t} + \mathbf{t}^T \mathbf{t} \quad (28)$$

$$= \mathbf{w}^T X^T X\mathbf{w} - 2\mathbf{t}^T X\mathbf{w} + \mathbf{t}^T \mathbf{t} \quad (29)$$

To optimize we aim to find the minimum, thus we differentiate with respect to \mathbf{w} using the identities from Table 1.4 (Rogers & Girolami, p.23):

$$\frac{\partial}{\partial \mathbf{w}} (\mathbf{w}^T X^T X\mathbf{w}) = 2X^T X\mathbf{w} \quad (30)$$

$$\frac{\partial}{\partial \mathbf{w}} (-2\mathbf{t}^T X\mathbf{w}) = -2X^T \mathbf{t} \quad (31)$$

$$\frac{\partial}{\partial \mathbf{w}} (\mathbf{t}^T \mathbf{t}) = \mathbf{0} \quad (32)$$

Thus, we find the gradient of the loss to be:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{w}} = 2X^T X\mathbf{w} - 2X^T \mathbf{t} \quad (33)$$

Setting the gradient to zero to find the minimum:

$$2X^T X\mathbf{w} - 2X^T \mathbf{t} = \mathbf{0} \quad (34)$$

$$X^T X\mathbf{w} = X^T \mathbf{t} \quad (35)$$

Assuming that $X^T X$ is invertible, we solve for \mathbf{w} :

$$(X^T X)^{-1} X^T X \mathbf{w} = (X^T X)^{-1} X^T \mathbf{t} \quad (36)$$

$$I\mathbf{w} = (X^T X)^{-1} X^T \mathbf{t} \quad (37)$$

Sine $Iw = w$ (identity matrix times an vector equals that vector) we have:

$$\hat{\mathbf{w}} = (X^T X)^{-1} X^T \mathbf{t} \quad (38)$$

If we look at the average loss equation from the literature (Rogers & Girolami, equation 1.13, p.21) we see that the equations differ by $\frac{1}{N}$:

$$L = \frac{1}{N}(\mathbf{t} - X\mathbf{w})^T(\mathbf{t} - X\mathbf{w}) \quad (39)$$

Following the same derivation as above we get:

$$\frac{\partial L}{\partial \mathbf{w}} = \frac{2}{N} X^T X \mathbf{w} - \frac{2}{N} X^T \mathbf{t} \quad (40)$$

Setting this to zero and solving:

$$\frac{2}{N} X^T X \mathbf{w} - \frac{2}{N} X^T \mathbf{t} = \mathbf{0} \quad (41)$$

$$X^T X \mathbf{w} = X^T \mathbf{t} \quad (42)$$

This yields the identical solution: $\hat{\mathbf{w}} = (X^T X)^{-1} X^T \mathbf{t}$. This happens as the constant factor $\frac{1}{N}$ in the average loss appears in every term of the gradient. When we set the gradient to zero and solve for w , the factor cancels out.

References

- [1] Simon Rogers and Mark Girolami. *A First Course in Machine Learning*. Machine Learning & Pattern Recognition. Chapman and Hall/CRC, Boca Raton, FL, 2 edition, 2016.

A housing_1.py

```
1 # load data
2 train_data = numpy.loadtxt("boston_train.csv", delimiter=",")
3 test_data = numpy.loadtxt("boston_test.csv", delimiter=",")
4 X_train, t_train = train_data[:, :-1], train_data[:, -1]
5 X_test, t_test = test_data[:, :-1], test_data[:, -1]
6 # make sure that we have N-dimensional Numpy arrays (ndarray)
7 t_train = t_train.reshape((len(t_train), 1))
8 t_test = t_test.reshape((len(t_test), 1))
9 print("Number of training instances: %i" % X_train.shape[0])
10 print("Number of test instances: %i" % X_test.shape[0])
11 print("Number of features: %i" % X_train.shape[1])
12
13 # (a) compute mean of prices on training set
14
15 mean = numpy.mean(t_train)
16 print(f"a) Mean house price on training set: ${mean:.2f}k")
17
18 # (b) RMSE function
19 def rmse(t, tp):
20     return numpy.sqrt(numpy.mean((t-tp)**2)) # begin by finding differences, square
21     ↳ the errors, and take the mean of the squared errors and finnaly take the square
22     ↳ root
23
24 # Create predictions (all equal to mean)
25 predictions = numpy.full_like(t_test, mean)
26
27 # Compute RMSE
28 test_rmse = rmse(t_test, predictions)
29 print(f"b) RMSE on test set: ${test_rmse:.2f}k")
30
31 # (c) visualization of results
32 plt.figure()
33 plt.scatter(t_test, predictions)
34 plt.xlabel("True price")
35 plt.ylabel("Predicted price")
36 plt.title("Mean Model")
37 plt.savefig("housing_1_scatter.png")
```

B housing_2.py

```
1 numpy.set_printoptions(precision=3) #for .3f for arrays
2
```

```
3 # load data
4 train_data = numpy.loadtxt("boston_train.csv", delimiter=",")
5 test_data = numpy.loadtxt("boston_test.csv", delimiter=",")
6 X_train, t_train = train_data[:, :-1], train_data[:, -1]
7 X_test, t_test = test_data[:, :-1], test_data[:, -1]
8 # make sure that we have N-dimensional Numpy arrays (ndarray)
9 t_train = t_train.reshape((len(t_train), 1))
10 t_test = t_test.reshape((len(t_test), 1))
11 print("Number of training instances: %i" % X_train.shape[0])
12 print("Number of test instances: %i" % X_test.shape[0])
13 print("Number of features: %i" % X_train.shape[1])
14
15 # (b) fit linear regression using only the first feature
16 model_single = linreg.LinearRegression()
17 model_single.fit(X_train[:, 0], t_train) # builds design matrix we = etc,
18
19 print(f"intercept: {model_single.w[0, 0]:.3f}")
20 print(f"CRIM: {model_single.w[1, 0]:.3f}")
21
22 # (c) fit linear regression model using all features
23 model_all = linreg.LinearRegression()
24 model_all.fit(X_train, t_train)
25
26 print(f"All weights:\n{model_all.w}")
27
28 # (d) evaluation of results
29
30 single_prediction = model_single.predict(X_test[:, 0])
31 all_prediction = model_all.predict(X_test)
32
33 #RMSE
34 single_rmse = numpy.sqrt(numpy.mean((single_prediction - t_test)**2))
35 all_rmse = numpy.sqrt(numpy.mean((all_prediction - t_test)**2))
36 print(f"RMSE single: {single_rmse:.3f}")
37 print(f"RMSE all: {all_rmse:.3f}")
38
39 plt.figure()
40 plt.scatter(t_test, single_prediction)
41 plt.xlabel("True price")
42 plt.ylabel("Predicted price")
43 plt.title("Single feature (CRIM)")
44 plt.savefig("scatter_single.png")
45
46 plt.figure()
47 plt.scatter(t_test, all_prediction)
48 plt.xlabel("True price")
49 plt.ylabel("Predicted price")
```

```
50 plt.title("All features")
51 plt.savefig("scatter_all.png")
52
53 plt.show()
54
```

C linreg.py

```
1 class LinearRegression:
2 """
3 Linear regression implementation.
4 """
5
6 def __init__(self):
7     pass
8
9
10 def fit(self, X, t):
11 """
12     Fits the linear regression model.
13
14     Parameters
15     -----
16     X : Array of shape [n_samples, n_features]
17     t : Array of shape [n_samples, 1]
18 """
19     if X.ndim == 1:
20         X = X.reshape(-1, 1) # converts 1d vector to matrix
21
22     ones = numpy.ones((X.shape[0], 1))
23     new_X = numpy.hstack([ones, X]) # make new matrix with 1s
24
25     # building the normal equation - Aw = b, solving for w
26     A = new_X.T @ new_X
27     b = new_X.T @ t
28     self.w = numpy.linalg.solve(A, b)
29
30 def predict(self, X):
31 """
32     Computes predictions for a new set of points.
33
34     Parameters
35     -----
36     X : Array of shape [n_samples, n_features]
37
```

```
38     Returns
39     -----
40     predictions : Array of shape [n_samples, 1]
41     """
42     if X.ndim == 1:
43         X = X.reshape(-1,1)
44     ones = numpy.ones((X.shape[0], 1))
45     new_X = numpy.hstack([ones, X])
46
47     return new_X @self.w
```